Final Paper, CS6460 (Educational Technology)

Chris Hodapp, `chodapp3@gatech.edu`

ABSTRACT

Many programming languages or environments have been created for the purpose of teaching programming interactively (such as Logo, Scratch, Processing, and Squeak), and have at their root Papert's notion of constructionism. At the same time, the web browser has become an increasingly common mode of delivery of technical tutorials and even textbook content, and accomodates increasingly sophisticated content to the point of that content not just including full graphical programs, but being a development platform in its own right. The project described here is an attempt to combine these in some fashion. More specifically, it presents an interactive 2D/3D programming and rendering environment via WebGL, uses this environment as a platform for content designed to teach some topics in mathematics, and delivers all of this via online content accessible from most contemporary web browsers. It is presently hosted online at: `https://hodapp87.github.io/cs6460_project`

1. INTRODUCTION

Computer programming has been used as a platform for teaching various subjects as computers became more and more widespread. The most notable example of this is perhaps via Seymour Papert's work in constructionism and his co-creation of the Logo programming language for teaching mathematics and analytical thought, primarily to children(Stager, 2016; Papert, 1993, 1980, 1999). His approach here stresses the importance in education of the learner being empowered to build things and being able to "see" from the perspective of some abstraction (such as the turtle in Logo) as a way of guiding and honing intuition.

Resnick, reflecting on the design of the Scratch language and its roots in Logo, also offers a succint summary (Resnick, Silverman, Kafai, Maloney, Monroy-Hernández, Rusk, Eastmond, Brennan, Millner, Rosenbaum, and Silver, 2009) of the balancing act of creating a programming language whose goal is to educate:

> Papert argued that programming languages should have a "low floor" (easy to get started) and a "high ceiling" (opportunities to create increasingly complex projects over time). In addition, languages need "wide walls" (supporting many different types of projects so people with many different interests and learning styles can all become engaged). Satisfying the triplet of low-floor/high-ceiling/wide-walls hasn't been easy.

The same author also gives a reason why one should use programming at all in teaching:

> The ability to program provides important benefits. For example, it greatly expands the range of what you can create (and how you can express yourself) with the computer. It also expands the range of what you can learn. In particular, programming supports "computational thinking," helping you learn important problem-solving and design strategies (such as modularization and iterative design) that carry over to nonprogramming domains. And since programming involves the creation of external representations of your problem-solving processes, programming provides you with opportunities to reflect on your own thinking, even to think about thinking itself.

In a somewhat related way, web content has made up a substantial amount of learning material, including at the same scope as a textbook. Some more-established examples of this are SICP (Structure and Interpretation of Computer Programs)(Abelson, Sussman, and Sussman, 1997), HtDP (How to Design Programs)(Felleisen, Findler, Flatt, and Krishnamurthi, 2003), and Software Foundations(Pierce, de Amorim, Casinghino, Gaboardi, Greenberg, Hriţcu, Sjöberg, and Yorgey, 2017). Overall, though, these are still something like "traditional" textbook content distributed in a way that takes advantage of electronic media in general, but does not make

any particular use of the nuances of web content or the web browser. Software Foundations, which teaches principles of computation and formal proofs using the Coq programming language, applies a notable method here: the learner can read the entire book in the Coq environment, and examples and exercises are provided in the form of Coq code which can be edited, compiled, and evaluated directly in that environment. SICP and HtDP have both been used in introductory computer science courses at colleges, including "traditional", non-online courses(Felleisen, Findler, Flatt, and Krishnamurthi, [n. d.]; Harvey, 2011).

At the same time, web browsers have also progressed substantially from their original purpose of accessing content. At the present, the web browser is a platform unto itself for running applications (via JavaScript), and these applications can be quite sophisticated: they may be written in many other programming languages (provided the language can compile to JavaScript), they may be highly visual and interactive with the help of things like Canvas and SVG and use hardware-accelerated 3D graphics by way of WebGL, these applications may even be development environments unto themselves (due to JavaScript's interpreted nature), and presently the majority of browsers across both desktop computers and mobile devices support all of these features simply as standard.

Learning material has made some good use of this; for instance, Eigenvectors and Eigenvalues Explained Visually and Seeing Theory: A visual introduction to probability and statistics made use of the visual capabilities of web browsers to present topics in mathematics. While this is not anything new per se (similar things are numerous in the form of desktop applications, Java applets, Flash applications, and mobile applications), these have the advantage that they can run nearly anywhere a modern web browser does and typically without requiring any special configuration, plugins, or other software.

For interactive learning materials for programming education, one common paradigm is that the web browser is used as an interface for resources that are mostly hosted server-side (typically, an implementation of some programming language), as technologies like Ajax made more feasible around 2005(Garrett, 2005). For instance, with the Go programming language interactive tutorial, the user edits the Go program in the browser, and the interface then submits the code to a server which is responsible for compiling it, executing it, and submitting the results back to the browser. While this is a very flexible method, it has drawbacks:

- It requires that the server have the right software, configuration, and resources to host the environment in question.

- It may require a reasonably fast Internet connection, or at least a low-latency one, for the interface to be usable. For any larger amount of data that the code is to handle - such as images, audio, or video - this can be problematic.

- Users are given the freedom to run arbitrary programs on a remote server, and the server must be equipped to prevent accidental or malicious misuse from causing problems. For instance, a user should not be able to delete or even access arbitrary files on the server, access privileged system calls, use all available memory or disk space, or indefinitely run a program that occupies a large amount of resources. This ranges between difficult and impossible to enforce at the language level, and as a result servers will commonly employ some kind of virtualization or other sandboxing to try to achieve this.

However, as mentioned, the web browser is able to run JavaScript code, and this ability has become increasingly sophisticated. One obvious use of this is that content can allow the learner to edit, debug, and execute JavaScript code directly in their browser (and indeed most browsers come built-in with development tools to this end) - such as JSFiddle. This has various benefits: the sandboxed nature of the browser prevents the learner from doing any real damage, the fact that the code only executes on one's own computer limits that damage further (as the sandboxing is never perfect), and any language implementation that compiles to JavaScript can be run in the browser. Tools such as Emscripten from Mozilla extend the reach of this considerably.

As an example of this, the Lean Theorem Prover has several tutorials (such as (Avigad, de Moura, Ebner, and Ullrich, 2017)) which are implemented in much the same way as Software Foundations, but hosted completely in the web browser. The fact that they could compile the implementation of the Lean language itself to JavaScript (that is, its compiler and runtime run directly in the browser) enabled them to host in this way without requiring any particular server-side resources. However, the user is not required to run this way; Lean is still available as a natively-compiled application or as source code.

The visual programming language Processing - itself designed for teaching programming visually(Processing, 2015) - provides another example of this with in the form of Processing.js, a version of Processing ported to JavaScript.

Programming using tooling outside of the web browser (e.g. in a dedicated IDE, using a text editor and some commandline tools, or at a REPL) is still a superior option for many contexts, particularly for programming tasks that are lower-level, more performance-critical, or more complex. However, these almost always come with some sort of barrier to entry that is in addition to the difficulty of learning the language in the first place: the tooling must be installed, sometimes licenses must be acquired, the tooling often has an interface with its own nuances that must be learned, and the tooling may require separate installation or management of plugins, libraries, packages, and so on.

These approaches are not necessarily exclusive of each other though. Environments based in the web browser can provide a sort of "low floor" here (as in the Papert quote), but can achieve the "high ceiling" and "wide

walls" part of things by making it more feasible for the user to transition to more dedicated tools when needed.

2. GOALS

In light of the current context of web browsers and their capabilities, the goal in this project was to create educational content, and an interactive programming environment, along the following lines:

- The content is hosted inside of an interactive programming environment which can run from a web browser, ideally with minimal server requirements (i.e.

- Content and implementation should be freely available as open source.

- The interactive programming environment is sufficient to allow the user most of the "normal" amenities in a language, includes built-in functionality for creating graphics, visualizations, and animation, and is advanced enough to accomodate some sophistication here (for instance, a raytracer or other 3D renderer).

- Content is directed at letting the learner explore topics in analytic geometry and linear algebra by way of writing programs which create graphics.

The content's area is something of a subtlety: It is not meant as introductory programming lessons, nor as a tutorial in how to create graphics or how to write a renderer, nor as solely mathematical content. It is meant as an explanation of some mathematical topics, using interactive programming of graphics as the method.

This was based partly around Papert's notion of making math "appropriable": "In many ways mathematics - for example the mathematics of space and movement and repetitive patterns of action - is what comes most naturally to children."(Papert, 1993) While the content isn't meant for children, but for younger adults with some existing programming and mathematics experience, the principle should hold similarly.

3. IMPLEMENTATION & CONTENT

Everything discussed below is available to use at `https://hodapp87.github.io/cs6460_project` and the source code for this is at `https://github.com/hodapp87/cs6460_project`.

3.1 Environment

This is based on a modified version of mkleanbook, an open source tool which is also the tool used to generate the Introduction to Lean mentioned earlier (and other interactive Lean tutorials).

The modifications were primarily to replace the tool's support for the Lean language with support more specific to WebGL. WebGL support was mostly done by integrating a modified version of pocket.gl which is available at `https://github.com/Hodapp87/pocket.gl`; this component was responsible for allowing the user to edit and render WebGL code.

More specifically, the project includes a build process which turns Emacs org-mode markup into a webpage in the form of a directory of static HTML/JavaScript/CSS that is suitable for hosting directly from a web server. This webpage contains content produced from that org-mode markup, but it also contains an interactive WebGL editor and renderer. This org-mode markup can contain segments of WebGL source code, and the build turns these into properly-highlighted source code in the rendered page, placing a "Try code" button below each segment which will transfer it directly to the editor.

See figure 1 for a screenshot showing an example of this (from the content described in the next section) with the Fragment Shader tab visible, and 2 for the same but with the Render tab showing the animated result of some code (along with its parameters).

Since everything was written to rely only on static files (and not requiring any special server-side support), this page could be hosted directly from its GitHub repository via GitHub Pages. It does not rely on anything specific to git or GitHub.

3.2 Content

The actual content of this project is done through the above environment. As planned, it tries to explain a small selection of topics in mathematics using examples in WebGL which the learner may execute and edit directly.

All code examples are implemented as standalone WebGL fragment shaders. That is, while the code can be hardware-accelerated, it uses the GPU (if there is one) almost solely as a stream processor rather than as dedicated rendering hardware, and so rather than focusing on OpenGL's rendering model, it focuses on what can be done in creating graphics pixel-by-pixel.

The main goal is in explaining and building intuition for the mathematical topics in question. A secondary goal is in empowering the learner to still be working with "real" code and comprehending its fundamental abstractions in such a way that it could easily be applied in other places - such as a larger OpenGL application, a homemade renderer, or any other shading language.

It is split into two main parts: a chapter with 2D graphics, and one with 3D graphics. The chapter on 2D graphics begins with creating graphics using so-called functional images(Elliott, 2003) - that is, an image that is produced by a function f which maps an (x, y) pixel location - a point in 2D space - to an RGB triplet, i.e. $[0, 1]^3$, to specify the color of pixel (x, y). The use of GL fragment shaders is not incidental here: fragment shaders work exactly according to this model.

The chapter walks through some mechanisms for working with the coordinate space and drawing basic shapes under this model, and then demonstrates how transforms applied to the domain of f are intuitively a sort of warping of space that can be comprehended visually, and how function composition relates to this. It also explains some built-in functionality for being able to both animate these graphics and interactively control parameters to them.
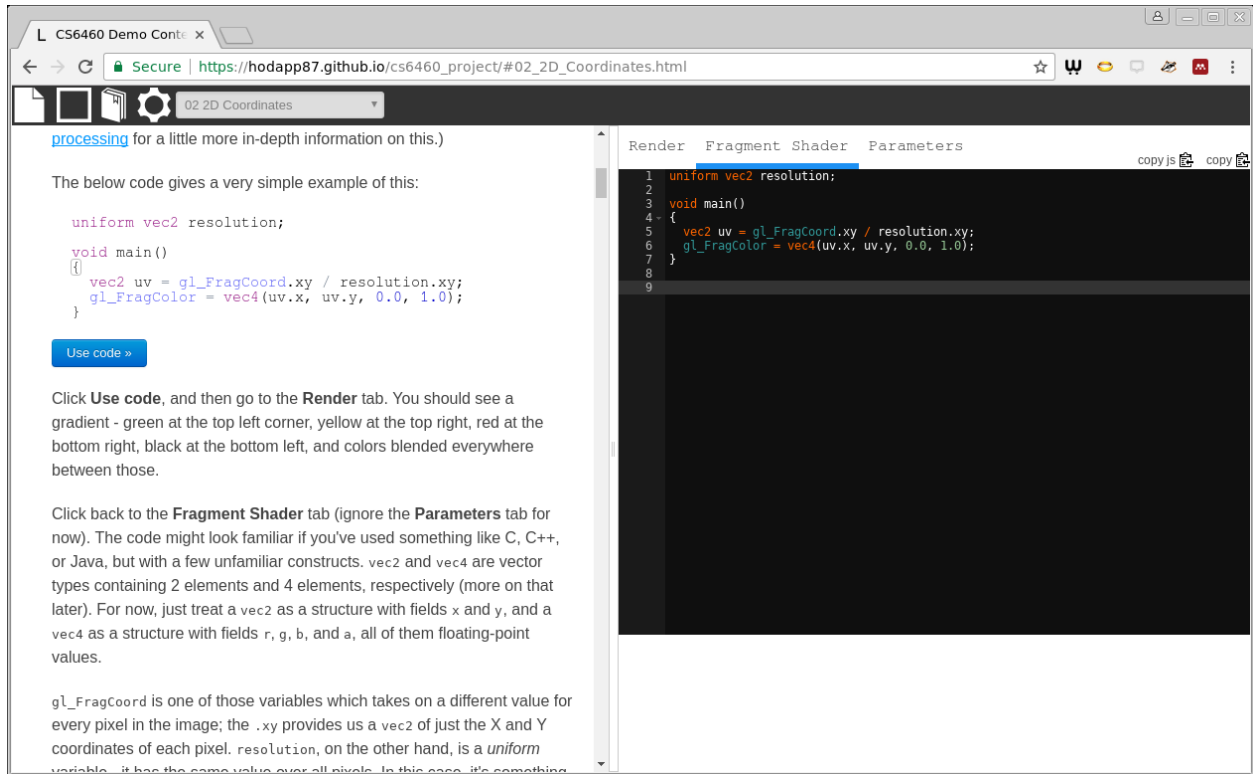
Secure | https://hodapp87.github.io/cs6460_project/#02_2D_Coordinates.html

02 2D Coordinates

Render   Fragment Shader   Parameters

copy js   copy

```
uniform vec2 resolution;

void main()
{
    vec2 uv = gl_FragCoord.xy / resolution.xy;
    gl_FragColor = vec4(uv.x, uv.y, 0.0, 1.0);
}
```

processing for a little more in-depth information on this.)

The below code gives a very simple example of this:

```
uniform vec2 resolution;

void main()
{
    vec2 uv = gl_FragCoord.xy / resolution.xy;
    gl_FragColor = vec4(uv.x, uv.y, 0.0, 1.0);
}
```

Use code »

Click **Use code**, and then go to the **Render** tab. You should see a gradient - green at the top left corner, yellow at the top right, red at the bottom right, black at the bottom left, and colors blended everywhere between those.

Click back to the **Fragment Shader** tab (ignore the **Parameters** tab for now). The code might look familiar if you've used something like C, C++, or Java, but with a few unfamiliar constructs. vec2 and vec4 are vector types containing 2 elements and 4 elements, respectively (more on that later). For now, just treat a vec2 as a structure with fields x and y, and a vec4 as a structure with fields r, g, b, and a, all of them floating-point values.

gl_FragCoord is one of those variables which takes on a different value for every pixel in the image; the .xy provides us a vec2 of just the X and Y coordinates of each pixel. resolution, on the other hand, is a *uniform* variable - it has the same value over all pixels. In this case, it's something

Figure 1: **Project web page with Fragment Shader visible**

Render   Fragment Shader   Parameters

Freq1   0.1
Freq2   -0.19
Rad   -0.04
Inv   0
SinAmp   5
SinFreq   -0.22
SinFreq2   1

Close Controls

Coordinates & Transforms in 2D

2 Coordinates & Transforms in 2D

2.1 A slightly boring, laborious example

The Introduction mentioned that all the code this tutorial deals with is *fragment shaders* (or pixel shaders) in GLSL. For our sake, what that means is that we are dealing with short programs which create images (and eventually animations) by explicitly giving the color of every pixel in the image.

This may be a bit different than what you are used to in more general programming languages. You are likely used to programs for which, at any given point in time, variables have a specific value and the program is executing some specific part of the code.

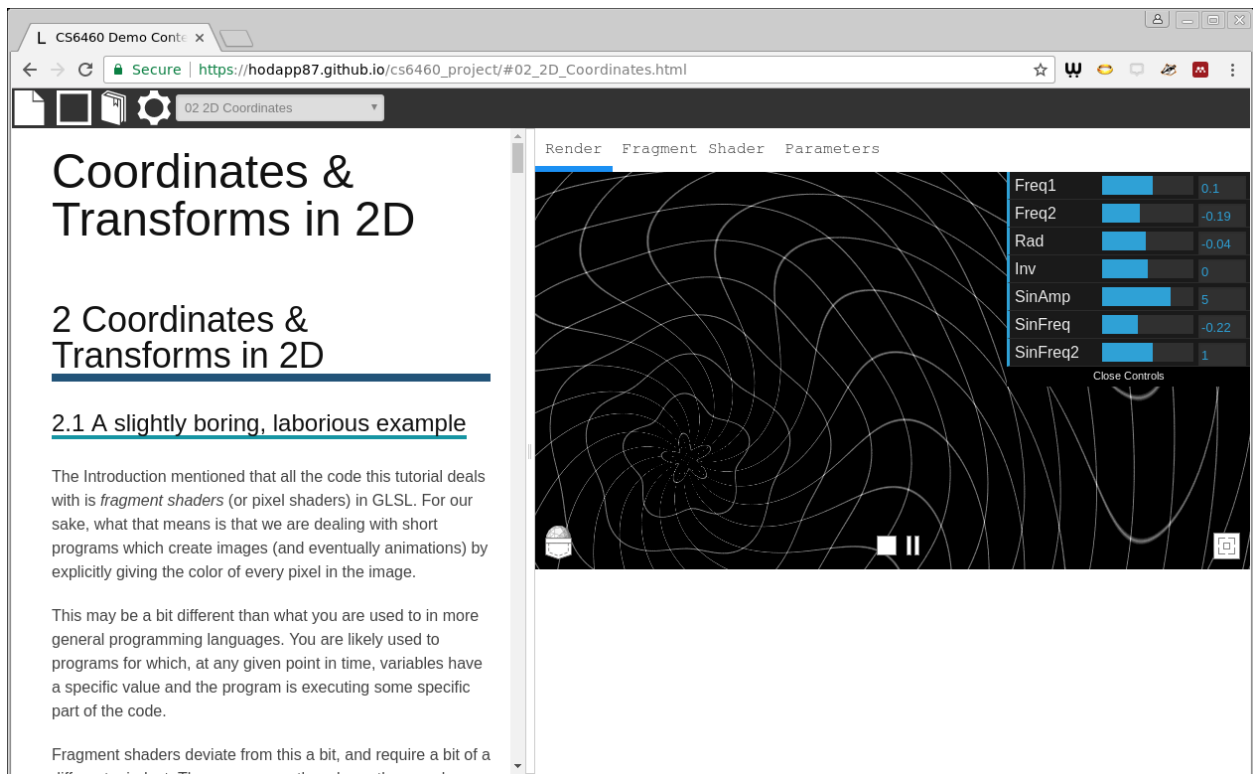Fragment shaders deviate from this a bit, and require a bit of a

Figure 2: **Project web page with Render visible**

3.2.1 Sphere Tracing

The next chapter presents one extension of this "functional images" model from 2D to 3D. In specific, rather than a functions which describes an image by mapping a 2D point (x, y) to an RGB triplet, it introduces distance functions which describe a 3D surface by mapping a 3D point (x, y, z) to the nearest distance from (x, y, z) to that surface (in other words, the radius r of the largest sphere located at (x, y, z) which never intersects the surface).

One use of these distance functions is that raymarching can be used to render them, and here a specific form of raymarching - sphere tracing - is used (Hart, 1994; Hart, Sandin, and Kauffman, 1989; Quilez, 2008; Perlin and Hoffert, 1989; Ebert, Musgrave, Peachey, Perlin, and Worley, 1998). An existing implementation of a sphere tracer from Íñigo Quílez is adapted for use in the example code.

Ray tracing is very commonly used in introducing a 3D renderer due to the simplicity of implementation, however, sphere tracing is used instead here because distance functions are much more flexible to the sorts of transformations and deformations that are applied. Figure 3 gives one simple example of a common effect done in distance functions in a sphere tracer: CSG is used to create a composite shape by subtracting a sphere from a cube, modulo is used to tile the shape infinitely in all directions, and a variable rotation is applied to this infinite field.

Like the prior chapter, this distance function representation is used to express some basic shapes. This is then used as the basis to express some rigid body 3D transformations via matrices and then via quaternions; this is then generalized to the same sorts of domain transformations that were done in the 2D case.

4. LIMITATIONS, PROBLEMS & FUTURE WORK

This work is considered to be some mix of a draft and a proof-of-concept. Besides minor bug-fixes and features in the environment itself, it has several areas for further improvement.

The content covered a limited scope of mathematical topics, mostly due to time constraints. It could, and perhaps should, extend to further topics that can be explained similarly. To name a few that were planned but not completed, the content was supposed to also include eigendecompositions, matrix inverses, fractals, and iterated function systems. The entire fourth chapter (at least, at the time of this writing) of the content contains a list of references to further topics which might sensibly be integrated in. The content also presently doesn't have the benefit of much feedback from people who attempted to learn from it.

Also due to time constraints, parts of the environment itself were adapted from existing libraries rather than being written from scratch. Some of these might more sensibly be rewritten from scratch (or at least something lower-lever). In somewhat the opposite way, some other parts that were written by hand might be better off

replaced with a separate module. To name one, the project could perhaps employ glslify to assist with the very ad-hoc management of packages presently in used in the environment's WebGL code.

The environment is also somewhat coupled to WebGL right now. Making the WebGL support more of a separate module may make this more useful to extend to other languages that might run in the browser.

Features were planned which would give the environment more of a "social" nature, such as the one provided in Scratch. The authors of Scratch distilled its design philosophy down to, "Make it more tinkerable, more meaningful, and more social than other programming environments."(Resnick et al., 2009) and spoke extensively of the importance of social networks in education. None of these features were implemented, however, environments like Scratch, ShaderToy, and OpenProcessing give some ideas for what form these might take.

References

Harold Abelson, Gerald Jay Sussman, and Julie Sussman. 1997. *Structure and Interpretation of Computer Programs*. MIT Press. `https://mitpress.mit.edu/sicp/full-text/book/book.html`

Jeremy Avigad, Leonarda de Moura, Gabriela Ebner, and Sebastian Ullrich. 2017. *An Introduction to Lean*. `https://leanprover.github.io/introduction_to_lean/`

David S. Ebert, F. Kenton Musgrave, Darwyn Peachey, Ken Perlin, and Steven Worley. 1998. *Texturing & Modeling: A Procedural Approach*. AP Professional.

Conal Elliott. 2003. Functional Images. In *The Fun of Programming*. Palgrave. `http://conal.net/papers/functional-images/`

Matthias Felleisen, Robert Bruce Findler, Matthew Flatt, and Shriram Krishnamurthi. [n. d.]. The Structure and Interpretation of the Computer Science Curriculum. ([n. d.]). `http://www.ccs.neu.edu/racket/pubs/jfp2004-fffk.pdf`

Matthias Felleisen, Robert Bruce Findler, Matthew Flatt, and Shriram Krishnamurthi. 2003. *How to Design Programs*. MIT Press. `http://www.htdp.org/`

Jesse James Garrett. 2005. Ajax: A New Approach to Web Applications. (2005). `http://adaptivepath.org/ideas/ajax-new-approach-web-applications/`

John C Hart. 1994. Sphere Tracing: A Geometric Method for the Antialiased Ray Tracing of Implicit Surfaces. (1994). `https://pdfs.semanticscholar.org/1f0d/001a247dc5b7facfa7ed445cc9d6a7801647.pdf`

John C. Hart, Daniel J. Sandin, and Louis H. Kauffman. 1989. Ray Tracing Deterministic 3-D Fractals. *Computer Graphics* 23, 3 (1989), 289–296. `https://doi.org/10.1145/74334.74363`

Brian Harvey. 2011. Why Structure and Interpretation of Computer Programs matters. (2011). `https://people.eecs.berkeley.edu/~bh/sicp.html`

Seymour Papert. 1980. Constructionism vs. Instruction-

**Figure 3: Render from sphere tracer used in project**

ism. (1980). http://www.papert.org/articles/const_inst/const_inst1.html

Seymour Papert. 1993. *Mindstorms: Children, Computers, and Powerful Ideas* (2 ed.). Basic Books.

Seymour Papert. 1999. What is Logo? Who Needs It? In *Logo Philosophy and Implementation.* Logo Computer Systems Inc., V–XVI. http://www.microworlds.com/company/philosophy.pdf

K. Perlin and E. M. Hoffert. 1989. Hypertexture. *ACM SIGGRAPH Computer Graphics* 23, 3 (1989), 253–262. https://doi.org/10.1145/74334.74359

Benjamin C. Pierce, Arthur Azevedo de Amorim, Chris Casinghino, Marco Gaboardi, Michael Greenberg, Cătălin Hriţcu, Vilhelm Sjöberg, and Brent Yorgey. 2017. *Software Foundations.* https://softwarefoundations.cis.upenn.edu/current/index.html

Processing 2015. Design Principles for p5.js. (2015). https://github.com/processing/p5.js/wiki/Design-Principles

Iñigo Quilez. 2008. Rendering Worlds with Two Triangles with raytracing on the GPU in 4096 bytes. (2008). http://iquilezles.org/www/material/nvscene2008/rwwtt.pdf

Mitchel Resnick, Brian Silverman, Yasmin Kafai, John Maloney, Andrés Monroy-Hernández, Natalie Rusk, Evelyn Eastmond, Karen Brennan, Amon Millner, Eric Rosenbaum, and Jay Silver. 2009. Scratch: Programming for All. *Commun. ACM* 52, 11 (2009), 60. https://doi.org/10.1145/1592761.1592779 arXiv:-

Gary S. Stager. 2016. Seymour Papert (1928âĂŞ2016). *Nature* 537, 7620 (sep 2016), 308–308. https://doi.org/10.1038/537308a